

# Ana Marasović, Curriculum Vitae, September 2025

[www.anamarasovic.com](http://www.anamarasovic.com)  
[ana.marasovic@utah.edu](mailto:ana.marasovic@utah.edu)

Merrill Engineering, Office 2166  
50 Central Campus Dr  
Salt Lake City, UT 84112

## RESEARCH INTERESTS

Natural Language Processing, Human-Centered AI, Interpretability

## EMPLOYMENT

**Assistant Professor, University of Utah, Kahelrt School of Computing**  
**Young Investigator, Allen Institute for AI**  
Courtesy appointment at the Paul G. Allen School of CSE, University of Washington

Jul 2022 – present  
Mar 2019 – Jun 2022

## EDUCATION

**Ph.D. in Computational Linguistics, Heidelberg University**  
*Honors: Magna cum laude*

**M.Sc. in Applied Mathematics, University of Zagreb**  
*Honors: Summa cum laude*

**B.Sc. in Mathematics, University of Zagreb**

2015 – 2019  
2013 – 2015  
2010 – 2013

## HONORS, AWARDS, AND FELLOWSHIPS

- EMNLP 2025 Outstanding Paper Candidate
- CoLM 2025 Spotlight Paper
- University of Utah One-U Responsible Artificial Intelligence Initiative (RAI) Faculty Fellow
- ACL 2023 Best Paper Award
- ACL 2023 Best Senior Area Chair Award
- SoCal NLP Symposium 2022 Best Paper Award
- ACL 2020 Best Paper Honorable Mention Award
- UC Berkeley EECS 2020 Rising Star
- EMNLP 2020 Outstanding Reviewer Award
- EMNLP 2019 Outstanding Reviewer Award

## FUNDED PROPOSALS

**Martian, Research Gift, \$131,000**  
Title: “Mechanistic Interpretability for AI Safety”

**Google, Research Gift, \$30,000**  
Title: “Evaluating LLMs as Web Assistants”

**Allen Institute for AI, Research Gift, \$100,000**  
Title: “Expressing Uncertainty During Explanation Generation”

Apr 2025  
Oct 2024  
Oct 2022

## PUBLICATIONS

Under Review

- [1] Nathan Stringham, Fateme Hashemi Chaleshtori, Xinyuan Yan, Zhichao Xu, Bei Wang, and **Ana Marasović**. *The Errors We Make Teaching About AI's Errors*.

## Papers in Rigorously Reviewed Journals & Conferences

- [1] Martin Tutek, Fateme Hashemi Chaleshtori, **Ana Marasović**, and Yonatan Belinkov. *Measuring Faithfulness of Chains of Thought by Unlearning Reasoning Steps*. EMNLP 2025. **Under review for Outstanding Paper Award**.
- [2] Alexander Gill, Abhilasha Ravichander, and **Ana Marasović**. *What Has Been Lost with Synthetic Evaluation?* Findings of EMNLP 2025.
- [3] Michael Clemens and **Ana Marasović**. *MixAssist: An Audio-Language Dataset for Co-Creative AI Assistance in Music Mixing*. CoLM 2025. **Oral Spotlight; one of 24 out of 418 accepted papers**.
- [4] Jesse Woo\*, Fateme Hashemi Chaleshtori\*, **Ana Marasović\*\***, and Kenneth Marino\*\*. *BriefMe: A Legal NLP Benchmark for Assisting with Legal Briefs*. Findings of ACL 2025.
- [5] Fateme Hashemi Chaleshtori, Atreya Ghosal, Alexander Gill, Purbid Bamroo, and **Ana Marasović**. *On Evaluating Explanation Utility for Human-AI Decision-Making in NLP*. Findings of EMNLP 2024.
- [6] Oliver Bentham, Nathan Stringham, and **Ana Marasović**. *Chain-of-Thought Unfaithfulness as Disguised Accuracy*. TMLR 2024.
- [7] Ashim Gupta, Rishanth Rajendhran, Nathan Stringham, Vivek Srikumar, and **Ana Marasović**. *Whispers of Doubt Amidst Echoes of Triumph in NLP Robustness*. NAACL 2024.
- [8] Jack Hessel, **Ana Marasović**, Jena D. Hwang, Lillian Lee, Jeff Da, Rowan Zellers, Robert Mankoff, and Yejin Choi. *Do Androids Laugh at Electric Sheep? Humor "Understanding" Benchmarks from The New Yorker Caption Contest*. ACL 2023. **Received a Best Paper Award given to only 4 out of 1074 accepted papers**. Covered in Psychology Today, The Debrief, and Cornell Chronicle.
- [9] Abhilasha Ravichander, Matt Gardner, and **Ana Marasović**. *CondaQA: A Contrastive Reading Comprehension Dataset for Reasoning about Negation*. EMNLP 2022. **Received a Best Paper Award at the SoCal NLP Symposium 2022; one of 3 award-winning papers out of 88 accepted papers**.
- [10] Alexis Ross, Matthew E. Peters, and **Ana Marasović**. *Does Self-Rationalization Improve Robustness to Spurious Correlations?* EMNLP 2022.
- [11] Shruti Palaskar, Akshita Bhagia, Yonatan Bisk, Florian Metze, Alan W Black, and **Ana Marasović**. *On Advances in Text Generation from Images Beyond Captioning: A Case Study in Self-Rationalization*. Findings of EMNLP 2022.
- [12] **Ana Marasović\***, Iz Beltagy\*, Doug Downey, and Matthew E. Peters. *Few-Shot Self-Rationalization with Natural Language Prompts*. Findings of NAACL 2022. **\*Equal contributions**.
- [13] Sarah Wiegrefe\* and **Ana Marasović\***. *Teach Me to Explain: A Review of Datasets for Explainable NLP*. NeurIPS Datasets and Benchmarks Track 2021. **\*Equal contributions**.
- [14] Sarah Wiegrefe, **Ana Marasović**, and Noah A. Smith. *Measuring Association Between Labels and Free-Text Rationales*. EMNLP 2021.
- [15] Jesse Dodge, Maarten Sap, **Ana Marasović**, William Agnew, Gabriel Ilharco, Dirk Groeneveld, Margaret Mitchell, and Matt Gardner. *Documenting Large Webtext Corpora: A Case Study on the Colossal Clean Crawled Corpus*. EMNLP 2021.
- [16] Kaiser Sun and **Ana Marasović**. *Effective Attention Sheds Light On Interpretability*. Findings of ACL 2021.
- [17] Alexis Ross, **Ana Marasović**, and Matthew E. Peters. *Explaining NLP Models via Minimal Contrastive Editing (MiCE)*. Findings of ACL 2021.
- [18] Alexander M. Hoyle, **Ana Marasović**, and Noah A. Smith. *Promoting Graph Awareness in Linearized Graph-to-Text Generation*. Findings of ACL 2021.

- [19] Alon Jacovi, **Ana Marasović**, Tim Miller, and Yoav Godlberg. *Formalizing Trust in Artificial Intelligence: Prerequisites, Causes and Goals of Human Trust in AI*. ACM FAccT 2021.
- [20] **Ana Marasović**, Chandra Bhagavatula, Jae sung Park, Ronan Le Bras, Noah A. Smith, and Yejin Choi. *Natural Language Rationales with Full-Stack Visual Reasoning: From Pixels to Semantic Frames to Commonsense Graphs*. Findings of EMNLP 2020.
- [21] Qiang Ning, Hao Wu, Pradeep Dasigi, Dheeru Dua, Matt Gardner, Robert L. Logan IV, **Ana Marasović**, and Zhen Nie. *Easy, Reproducible and Quality-Controlled Data Collection with CROWDAQ*. EMNLP System Demonstrations Track 2020.
- [22] Suchin Gururangan, **Ana Marasović**, Swabha Swayamdipta, Kyle Lo, Iz Beltagy, Doug Downey, and Noah A. Smith. *Don't Stop Pretraining: Adapt Language Models to Domains and Tasks*. ACL 2020. **Received an Honorable Mention for Best Paper Award; one of 5 award-winning papers out of 779 accepted papers.**
- [23] Pradeep Dasigi, Nelson F. Liu, **Ana Marasović**, Noah A. Smith, and Matt Gardner. *Quoref: A Reading Comprehension Dataset with Questions Requiring Coreferential Reasoning*. EMNLP 2019.
- [24] **Ana Marasović** and Anette Frank. *SRL4ORL: Improving Opinion Role Labeling Using Multi-Task Learning with Semantic Role Labeling*. NAACL 2018.
- [25] **Ana Marasović**, Leo Born, Juri Opitz, and Anette Frank. *A Mention-Ranking Model for Abstract Anaphora Resolution*. EMNLP 2017.
- [26] **Ana Marasović**, Mengfei Zhou, Alexis Palmer, and Anette Frank. *Modal Sense Classification At Large: Paraphrase-Driven Sense Projection, Semantically Enriched Classification Models and Cross-Genre Evaluations*. Linguistic Issues in Language Technology (LiLT) – Modality: Logic, Semantics, Annotation, and Machine Learning 14 (3). 2016

## Workshop Papers & Other Published Work

- [1] Purbid Bambroo, Daniel S. Brown, and **Ana Marasović**. *ConeCut: Scalable Removal of Preference Redundancy*. CoLM 2025 Workshop on Socially Responsible Language Modelling Research (SoLaR).
- [2] Karolina Stańczak, Nicholas Meade, Mehar Bhatia, Hattie Zhou, Konstantin Böttinger, Jeremy Barnes, Jason Stanley, Jessica Montgomery, Richard Zemel, Nicolas Papernot, Nicolas Chapados, Denis Thérien, Timothy P. Lillicrap, **Ana Marasović**, Sylvie Delacroix, Gillian K. Hadfield, and Siva Reddy. *Societal Alignment Frameworks Can Improve LLM Alignment*. ICLR 2025 Workshop on Bidirectional Human-AI Alignment.
- [3] Fateme Hashemi Chaleshtori, Atreya Ghosal, and **Ana Marasović**. *On Evaluating Explanation Utility for Human-AI Decision-Making in NLP*. The workshop “XAI in Action: Past, Present, and Future Applications” co-located with NeurIPS 2023.
- [4] Jacob K. Johnson and **Ana Marasović**. *How Much Consistency Is Your Accuracy Worth?* The Workshop on Analyzing and Interpreting Neural Networks for NLP (BlackboxNLP) co-located with EMNLP 2023.
- [5] Markus Zopf, Teresa Botschen, Tobias Falke, Benjamin Heinzerling, **Ana Marasović**, Todor Mihaylov, S. Avinesh P.V., Eneldo Loza Mencía, J. Fürnkranz, and A. Frank. *What's Important in a Text? An Extensive Evaluation of Linguistic Annotations for Summarization*. SNAMS 2018.
- [6] **Ana Marasović** and Anette Frank. *Multilingual Modal Sense Classification using a Convolutional Neural Network*. Workshop on Representation Learning for NLP (Repl4NLP) 2016.

## Other

- Kenneth Marino and **Ana Marasović**. *Computer Use Survey: A Visual Survey of Computer Use Agents*. 2025
- **Ana Marasović**. *NLP's Generalization Problem, and How Researchers Are Tackling It*. The Gradient 2018.

## TALKS

---

### Upcoming Talks

- In-person Invited talk at the University of Arizona Cognitive Science Colloquium Spring 2026
- Keynote at the First Workshop on the Application of LLM Explainability to Reasoning and Planning Oct 2025

### Keynote Talks

- If You Want Reasoning, Look Inside** May 2025  
The 10th Workshop on Representation Learning for NLP (RepL4NLP 2025) co-located with NAACL 2025
- Challenges in Fostering (Dis)Trust in AI** Oct 2023  
Workshop on “Trust Perspectives in Machine Learning, Law, and Public Policy” organized by the Institute for Data, Econometrics, Algorithms, and Learning; located at Northwestern University
- Contrastive Explanations of NLP Models** Nov 2021  
EMNLP Workshop on “Analyzing and Interpreting Neural Networks for NLP” (BlackBoxNLP)
- Resolving Abstract Anaphors in Discourse: Uphill Battles with Neural Ranking Models and Automatic Data Extraction** Jun 2018  
NAACL Workshop on Computational Models of Reference, Anaphora, and Coreference (CRAC)

### Invited Talks

- Explainability and Trust** Apr 2024  
Bellairs Invitational Workshop on Contemporary, Foreseeable and Catastrophic Risks of LLMs; Barbados
- Challenges in Fostering (Dis)Trust in AI** Apr 2024  
University of Pennsylvania  
University of Pittsburgh
- “We propose <insert an explainability method> to increase human trust in AI”: What does this mean and how to go about it?** Apr 2023  
University of Maryland
- “AI” That Masters Language Could Reason about Negation** Feb 2023  
University of Utah Data Science Seminar
- Formalizing Trust in AI** May 2022  
NSF AI Institute for Research on Trustworthy AI in Weather, Climate, and Coastal Oceanography (AI2ES)
- Self-Explaining for Intuitive Interaction with AI**  
Carnegie Mellon University Apr 2022  
NEC Labs Europe Apr 2022  
University of Lisbon - Unbabel Seminar Apr 2022  
UC Irvine Mar 2022  
Johns Hopkins University Feb 2022  
Brown University Feb 2022  
University of Chicago Feb 2022  
Ohio State University Feb 2022  
Purdue University, CS Feb 2022  
University of Utah Feb 2022
- On Reducing Human-Authored Free-Text Explanations for Training** Dec 2021  
University College London
- Explanation Selection Through The Lens of Free-Text & Contrastive Explanations**  
USC Machine Learning Symposium Dec 2021

|                                                                                                                          |          |
|--------------------------------------------------------------------------------------------------------------------------|----------|
| Georgia Tech NLP Seminar                                                                                                 | Nov 2021 |
| University of Sheffield Seminar                                                                                          | Nov 2021 |
| Purdue University, ECE                                                                                                   | Nov 2021 |
| <b>Explaining with Natural Language</b>                                                                                  | Oct 2021 |
| University of Edinburgh                                                                                                  |          |
| <b>Reflecting on Interpretability Desiderata with Visual Commonsense Reasoning</b>                                       | Nov 2019 |
| University of Washington, Department of Linguistics Seminar                                                              |          |
| <b>Teaching Machines Language Understanding</b>                                                                          | Apr 2019 |
| Allen School Women's Research Day, University of Washington                                                              |          |
| <b>Resolving Abstract Anaphors in Discourse: Uphill Battles with Neural Ranking Models and Automatic Data Extraction</b> |          |
| Allen Institute for AI                                                                                                   | Oct 2018 |
| Potsdam University, Applied Computational Linguistics                                                                    | Jun 2018 |
| Bochum University, Sprachwissenschaftliches Institut                                                                     | May 2018 |

## TEACHING

---

|                                                                  |                             |
|------------------------------------------------------------------|-----------------------------|
| <b>CS 5340/6340: Natural Language Processing</b>                 | Fall 2025, Spring/Fall 2024 |
| <b>CS 7930: Mechanistic Interpretability Seminar</b>             | Spring 2025                 |
| <b>CS 7931: Advances in Generative AI Seminar</b>                | Fall 2024                   |
| <b>CS 5966/6966: Local Explanations for Deep Learning Models</b> | Fall 2023, 2022             |
| <b>CS 5140/6140: Data Mining</b>                                 | Spring 2023                 |

## GUEST LECTURES & CONFERENCE TUTORIALS & PRESENTATIONS AT LOCAL EVENTS

---

|                                                                                              |            |
|----------------------------------------------------------------------------------------------|------------|
| <b>NAACL'24 Tutorial "Explanation in the Era of Large Language Models"</b>                   | June 2024  |
| <b>Ingredients of Generative AI</b>                                                          |            |
| University of Utah, Research Challenges in Computing Seminar, <i>Host</i> : Saday Sadayappan | Nov 2023   |
| University of Utah, DELPHI Data Science Symposium                                            | Oct 2023   |
| University of Utah, The Osher Lifelong Learning Institute Lunch and Learn                    | Oct 2023   |
| <b>NLP with (Almost) No Training Data</b>                                                    |            |
| University of Utah, MATH 4100 / COMP 5360, <i>Host</i> : Kate Isaacs                         | March 2023 |
| University of Utah, CS 5340/6340, <i>Host</i> : Ellen Riloff                                 | Nov 2022   |
| <b>Towards Trustworthy NLP: On Intuitive Interaction and Rigorous Validation</b>             |            |
| University of Utah, CS 3020 Research Forum, <i>Host</i> : Kobus Van Der Merwe                | Sep 2022   |
| <b>A Quick Tour of NLP Explainability</b>                                                    |            |
| University of Washington, <i>Host</i> : Yulia Tsvetkov                                       | May 2022   |
| <b>Detecting Artifacts with Local Explanations</b>                                           |            |
| University of British Columbia, <i>Host</i> : Vered Shwartz                                  | March 2022 |
| University of Washington, <i>Host</i> : Zaid Harchaoui                                       | March 2022 |
| <b>Methods for Analyzing (Current) NLP Models</b>                                            |            |
| University of Washington, <i>Host</i> : Noah A. Smith                                        | March 2020 |

## MENTORING

---

### Utah PhD Students

|                                   |                    |
|-----------------------------------|--------------------|
| <b>Fateme Hashemi Chaleshtori</b> | Jan 2023 – present |
| <b>Nathan Stringham</b>           | Aug 2023 – present |
| <b>Alex Gill</b>                  | Aug 2025 – present |
| <b>Braden Fiedel</b>              | Aug 2025 – present |

### Utah PhD Student Collaborators

|                                                                 |                     |
|-----------------------------------------------------------------|---------------------|
| <b>Michael Clemens</b> → New Jersey Institute of Technology PhD | Jul 2023 – present  |
| <b>Zhichao Xu</b>                                               | Dec 2023 – present  |
| <b>Atreya Ghosal</b>                                            | Apr 2023 – Jun 2024 |
| <b>Oliver Benthham</b>                                          | Jan 2024 - Feb 2024 |
| <b>Ashim Gupta</b>                                              | Oct 2022 – Dec 2023 |
| <b>Akansha Kalra</b>                                            | Nov 2022 – Dec 2022 |

### Utah MS Researchers

|                                                         |                                     |
|---------------------------------------------------------|-------------------------------------|
| <b>Purbid Bambroo</b> → Noetica                         | Jan – May 2024; Sep 2024 – Aug 2025 |
| <b>Alex Gill</b> → University of Utah PhD               | Jan 2024 – Aug 2025                 |
| <b>Rishanth Rajendhran</b> → University of Maryland PhD | Nov 2022 – May 2024                 |
| <b>Adwait Shinganwade</b>                               | Aug 2024 – Dec 2024                 |
| <b>Yash Lele</b>                                        | Aug 2024 – Dec 2024                 |
| <b>Mugdha Abhyankar</b>                                 | Jan 2023 – May 2023                 |

### Utah BS Researchers

|                                                                            |                     |
|----------------------------------------------------------------------------|---------------------|
| <b>Braden Fiedel</b> → University of Utah PhD                              | Aug 2024 – Aug 2025 |
| <b>Casey Miller</b>                                                        | Aug 2024 – present  |
| <b>Kishan Thambu</b>                                                       | Aug 2024 – present  |
| <b>Byron Liu</b> (co-mentoring with the PhD student <b>Maitrey Mehta</b> ) | Jul 2023 – May 2024 |
| <b>Jordan Tan</b>                                                          | Jul 2023 – Aug 2023 |
| <b>Dana Ahn</b>                                                            | Mar 2023 – Aug 2023 |
| <b>Jacob Johnson</b> → Utah PhD → Weber State University MS                | Nov 2022 – May 2024 |

### UW BS/MS Researchers

|                                                                            |                     |
|----------------------------------------------------------------------------|---------------------|
| <b>Kaiser Sun</b> with Noah A. Smith → META Residency → JHU PhD            | Mar 2020 – present  |
| <b>Millicent Li</b> with Noah A. Smith → META Residency → Northeastern PhD | Mar 2020 – Aug 2021 |

### AI2 Predoctoral Researchers

|                                                     |                       |
|-----------------------------------------------------|-----------------------|
| <b>Alexis Ross</b> with Matthew E. Peters → MIT PhD | July 2020 – June 2022 |
|-----------------------------------------------------|-----------------------|

### AI2 Interns (PhD students)

|                                                                                           |                      |
|-------------------------------------------------------------------------------------------|----------------------|
| <b>Abhilasha Ravichander</b> – Carnegie Mellon University – with Matt Gardner             | Jun 2021 – June 2022 |
| <b>Shruti Palaskar</b> – Carnegie Mellon University                                       | May 2021 – June 2022 |
| <b>Mimansa Jaiswal</b> – University of Michigan                                           | Sep 2021 – Dec 2021  |
| <b>Alexander M. Hoyle</b> – University of Maryland – with Noah A. Smith                   | Jun 2020 – Nov 2020  |
| <b>Sarah Wiegrefe</b> – Georgia Tech – <b>AI2 Intern of the Year</b> – with Noah A. Smith | May 2020 – Oct 2020  |
| <b>Zhanming Allan Jie</b> – Singapore Uni. of Technology & Design – with Pradeep Dasigi   | Jun 2019 – Aug 2019  |

## SERVICE

---

### Department & University Service

|                                                           |                     |
|-----------------------------------------------------------|---------------------|
| Mentoring for Grad Students                               | Sep 2025 – present  |
| Faculty Upskilling in AI (U-AI) instructor and consultant | Fall 2024           |
| AI Track Director                                         | Jul 2023 – present  |
| One Utah RAI Healthcare Working Group                     | Feb 2024 – Jul 2024 |

Faculty Hiring Committee  
PhD (AI) Admissions Committee

Nov 2023 – May 2024  
Dec 2022 – Jan 2023

### **PhD Committee Member**

Laura Vietz  
Ashim Gupta  
Maitrey Mehta  
Jakob Johnson  
Mattia Medina-Grespan  
Zhichao Xu  
Peter Hase (UNC Chapel Hill)  
Zhenduo Wang  
Marcos Treviso (University of Lisbon)  
Prince Osei Aboagye  
Tianyu Jiang

### **Co-organizer**

TeachingNLP: Workshop on Teaching Natural Language Processing @ ACL 2024  
SustainNLP: Workshop on Simple and Efficient Natural Language Processing @ EMNLP 2021  
AllenNLP 2021 Hackathon for Minoritized Communities  
Host of the NLP Highlights Podcast from 2021-2022  
AIPHES Scientific Workshop 2017

### **Best Paper Award Committee**

EMNLP 2025  
NAACL 2024

### **Senior Area Chair**

ACL 2024  
ACL 2023 (**Best Senior Area Chair**)  
EMNLP 2022

### **Monograph Reviewer**

Reviewer for “Synthesis Lectures in Human Language Technologies: Explainable Natural Language Processing” in 2021

### **Journal Reviewer**

TACL 2017 (secondary reviewer), 2021; *in 2021 TACL had an impact factor of 9.194*

### **Conference Reviewer**

ARR May 2005, ARR July 2025  
NAACL 2024  
EMNLP 2018, 2019 (**outstanding reviewer**), 2020 (**outstanding reviewer**), 2021, 2024  
ACL 2017 (secondary reviewer), 2019, 2020, 2021  
EACL 2021, 2023  
ICML 2020  
AAAI 2020

### **Workshop Reviewer**

Representation Learning for NLP (Repl4NLP) @ ACL 2018, ACL 2019  
Deep Learning Approaches for Low-Resource Natural Language Processing (DeepLo) @ EMNLP 2019  
Learning from Limited Labeled Data (LLD) Workshop @ ICLR 2019  
Evaluation and Comparison of NLP Systems (EvalNLP) @ EMNLP 2020  
Insights from Negative Results in NLP @ EMNLP 2020  
Computational Approaches to Discourse (CODI) @ EMNLP 2020  
Domain Adaptation for NLP (Adapt-NLP) @ EACL 2021

### **Panelist**

Moderator of the panel “AI in Pop Culture vs. Reality” at the Leonardo in Nov 2024  
Panelist at the Utah Digital Health Initiative “Generative AI in Health Sciences Research” panel in Nov 2023  
Panelist at the SCI’s 30th Anniversary event in April 2024  
Panelist at a career panel at 2023 EACL Student Research Workshop (SRW)  
Panelist for the screening of the movie “After Yang” for the STEM Education Film Series by Park City Film in Sep 2022  
Panelist at ACL Mentorship: “Building NLP datasets” in Aug 2022  
Panelist at SoCal ML & NLP Symposium 2021 (Networking and career paths for under-represented minorities)  
Panelist at undergraduate panel session @ ACL 2020

### **Moderator at Conferences**

Moderator of a meetup on “Interpretability” @ ACL 2021  
Moderator of a meetup on “Whether to do a PhD; and how to apply for and choose a PhD program” @ EMNLP 2020  
Moderator of a meetup on “Interpretability and Analysis of Models for NLP” at EMNLP 2020  
Moderator of a roundtable on “Natural Language Generation” at WiNLP @ ACL 2020

### **Mentoring at Conferences**

Mentor at WiNLP @ ACL 2020  
“Big sibling” at EMNLP 2019, NAACL 2019

### **Other Service**

Leonardo AI Advisory Committee Member, Jun 2024 – present  
Co-Lead of the AI2 Employee Resource Group (ERG) for Women and Gender Minorities in 2022  
Member of the AI2 Diversity, Equity and Inclusion Committee in 2020 and 2021  
Reader of 2020 PhD applications at the University of Washington

## **OPEN-SOURCE SOFTWARE AND DATASETS**

---

- Few-Shot Self-Rationalization with Natural Language Prompts  
<https://github.com/allenai/feb>
- Natural Language Rationales with Full-Stack Visual Reasoning: From Pixels to Semantic Frames to Commonsense Graphs  
<https://github.com/allenai/visual-reasoning-rationalization>
- Don’t Stop Pretraining: Adapt Language Models to Domains and Tasks  
Contributor  
<https://github.com/allenai/dont-stop-pretraining>
- Quoref Dataset  
Contributor  
<https://allenai.org/data/quoref>
- SRL4ORL: Improving Opinion Role Labeling Using Multi-Task Learning With Semantic Role Labeling  
<https://github.com/amarasovic/naacl-mpqa-srl4orl>
- A Mention-Ranking Model for Abstract Anaphora Resolution  
<https://github.com/amarasovic/neural-abstract-anaphora>  
<https://github.com/amarasovic/abstract-anaphora-data>
- Multilingual Modal Sense Classification using a Convolutional Neural Network  
<https://github.com/amarasovic/modal-sense-classification>  
*Dataset:* <http://projects.cl.uni-heidelberg.de/modals/>
- Matplotlib Examples for Visualizing Paper Results  
<https://github.com/amarasovic/matplotlib4papers>

## **REFERENCES**

---

- Noah A. Smith, [nasmith@cs.washington.edu](mailto:nasmith@cs.washington.edu)
- Yejin Choi, [yejin@cs.washington.edu](mailto:yejin@cs.washington.edu)
- Anette Frank, [frank@cl.uni-heidelberg.de](mailto:frank@cl.uni-heidelberg.de)
- Doug Downey, [dougd@allenai.org](mailto:dougd@allenai.org)