

www.anamarasovic.com
ana.marasovic@utah.edu

Merrill Engineering, Office 2166
 50 Central Campus Dr
 Salt Lake City, UT 84112

Research Interests

Natural Language Processing, Human-Centered AI, Interpretability

Employment

Assistant Professor , University of Utah, Kahlert School of Computing	Jul 2022 – present
Young Investigator , Allen Institute for AI	Mar 2019 – Jun 2022
Courtesy appointment at the Paul G. Allen School of CSE, University of Washington	

Education

Ph.D. in Computational Linguistics , Heidelberg University	2015 – 2019
<i>Honors</i>: Magna cum laude	
M.Sc. in Applied Mathematics , University of Zagreb	2013 – 2015
<i>Honors</i>: Summa cum laude	
B.Sc. in Mathematics , University of Zagreb	2010 – 2013

Honors, Awards, and Fellowships

Price College of Engineering 2026 Outstanding Teaching Award
 Kahlert School of Computing 2025–2026 Outstanding Teaching Award
 EMNLP 2025 Outstanding Paper
 CoLM 2025 Spotlight Paper
 University of Utah One-U Responsible Artificial Intelligence Initiative (RAI) Faculty Fellow
 ACL 2023 Best Paper Award
 ACL 2023 Best Senior Area Chair Award
 SoCal NLP Symposium 2022 Best Paper Award
 ACL 2020 Best Paper Honorable Mention Award
 UC Berkeley EECS 2020 Rising Star
 EMNLP 2020 Outstanding Reviewer Award
 EMNLP 2019 Outstanding Reviewer Award

Funding

Coefficient Giving , <i>Grant</i> , \$205,356; PI, collaboration with PI Daniel Brown	Nov 2025
Title: “Mechanistic Understanding of the Alignment of Large, Transformer-based, Learned Reward Models”	
Martian , <i>Research Gift</i> , \$131,000	Apr 2025
Title: “Mechanistic Interpretability for AI Safety”	
Google , <i>Research Gift</i> , \$30,000	Oct 2024
Title: “Evaluating LLMs as Web Assistants”	
University of Utah One-U RAI Initiative , <i>Faculty Fellowship</i> , \$135,000	2025–2027
Allen Institute for AI , <i>Research Gift</i> , \$100,000	Oct 2022
Title: “Expressing Uncertainty During Explanation Generation”	

Preprints

- [1] *Faithfulness Metrics Don't Measure Faithfulness: A Meta-Evaluation with Ground Truth*
Yoav Gur-Arieh, **Ana Marasović**, Mor Geva

Papers in Rigorously Reviewed Journals & Conferences

- [1] *The Hard Part Comes After Search: Benchmarking Web Agents on Synthesizing, Organizing, and Displaying Knowledge*
Alexander Gill, Md Farhan Ishmam, Xuyen Nguyen, Neha Bhat, Parker Henry DeYoung, Fateme Hashemi Chaleshtori, Nathan Stringham, Kenneth Marino, **Ana Marasović** • Findings of EMNLP 2026
- [2] *Societal Alignment Frameworks Can Improve LLM Alignment*
Karolina Stańczak, Nicholas Meade, Mehar Bhatia, Hattie Zhou, Konstantin Böttinger, Jeremy Barnes, Jason Stanley, Jessica Montgomery, Richard Zemel, Nicolas Papernot, Nicolas Chapados, Denis Therien, Timothy P. Lillicrap, **Ana Marasović**, Sylvie Delacroix, Gillian K. Hadfield, Siva Reddy • ACM FAccT 2026
- [3] *Measuring Faithfulness of Chains of Thought by Unlearning Reasoning Steps*
Martin Tutek, Fateme Hashemi Chaleshtori, **Ana Marasović**, Yonatan Belinkov • EMNLP 2025 • **Outstanding Paper**
- [4] *What Has Been Lost with Synthetic Evaluation?*
Alexander Gill, Abhilasha Ravichander, **Ana Marasović** • Findings of EMNLP 2025
- [5] *MixAssist: An Audio-Language Dataset for Co-Creative AI Assistance in Music Mixing*
Michael Clemens, **Ana Marasović** • CoLM 2025 • **Oral Spotlight**
- [6] *BriefMe: A Legal NLP Benchmark for Assisting with Legal Briefs*
Jesse Woo*, Fateme Hashemi Chaleshtori*, **Ana Marasović**, Kenneth Marino • Findings of ACL 2025
- [7] *On Evaluating Explanation Utility for Human-AI Decision-Making in NLP*
Fateme Hashemi Chaleshtori, Atreya Ghosal, Alexander Gill, Purbid Bambroo, **Ana Marasović** • Findings of EMNLP 2024
- [8] *Chain-of-Thought Unfaithfulness as Disguised Accuracy*
Oliver Bentham, Nathan Stringham, **Ana Marasović** • TMLR 2024
- [9] *Whispers of Doubt Amidst Echoes of Triumph in NLP Robustness*
Ashim Gupta, Rishanth Rajendhran, Nathan Stringham, Vivek Srikumar, **Ana Marasović** • NAACL 2024
- [10] *Do Androids Laugh at Electric Sheep? Humor “Understanding” Benchmarks from The New Yorker Caption Contest*
Jack Hessel, **Ana Marasović**, Jena D. Hwang, Lillian Lee, Jeff Da, Rowan Zellers, Robert Mankoff, Yejin Choi • ACL 2023 • **Best Paper Award** • Covered in Psychology Today, The Debrief, and Cornell Chronicle.
- [11] *CondaQA: A Contrastive Reading Comprehension Dataset for Reasoning about Negation*
Abhilasha Ravichander, Matt Gardner, **Ana Marasović** • EMNLP 2022 • **Best Paper Award at the SoCal NLP Symposium 2022**
- [12] *Does Self-Rationalization Improve Robustness to Spurious Correlations?*
Alexis Ross, Matthew E. Peters, **Ana Marasović** • EMNLP 2022
- [13] *On Advances in Text Generation from Images Beyond Captioning: A Case Study in Self-Rationalization*
Shruti Palaskar, Akshita Bhagia, Yonatan Bisk, Florian Metze, Alan W Black, **Ana Marasović** • Findings of EMNLP 2022
- [14] *Few-Shot Self-Rationalization with Natural Language Prompts*
Ana Marasović*, Iz Beltagy*, Doug Downey, Matthew E. Peters • Findings of NAACL 2022
- [15] *Teach Me to Explain: A Review of Datasets for Explainable NLP*
Sarah Wiegrefe*, **Ana Marasović*** • NeurIPS Datasets and Benchmarks Track 2021
- [16] *Measuring Association Between Labels and Free-Text Rationales*
Sarah Wiegrefe, **Ana Marasović**, Noah A. Smith • EMNLP 2021
- [17] *Documenting Large Webtext Corpora: A Case Study on the Colossal Clean Crawled Corpus*
Jesse Dodge, Maarten Sap, **Ana Marasović**, William Agnew, Gabriel Ilharco, Dirk Groeneveld, Margaret Mitchell, Matt Gardner • EMNLP 2021
- [18] *Effective Attention Sheds Light On Interpretability*
Kaiser Sun, **Ana Marasović** • Findings of ACL 2021
- [19] *Explaining NLP Models via Minimal Contrastive Editing (MiCE)*
Alexis Ross, **Ana Marasović**, Matthew E. Peters • Findings of ACL 2021

- [20] *Promoting Graph Awareness in Linearized Graph-to-Text Generation*
Alexander M. Hoyle, **Ana Marasović**, Noah A. Smith • Findings of ACL 2021
- [21] *Formalizing Trust in Artificial Intelligence: Prerequisites, Causes and Goals of Human Trust in AI*
Alon Jacovi, **Ana Marasović**, Tim Miller, Yoav Goldberg • ACM FAccT 2021
- [22] *Natural Language Rationales with Full-Stack Visual Reasoning: From Pixels to Semantic Frames to Commonsense Graphs*
Ana Marasović, Chandra Bhagavatula, Jae sung Park, Ronan Le Bras, Noah A. Smith, Yejin Choi • Findings of EMNLP 2020
- [23] *Easy, Reproducible and Quality-Controlled Data Collection with CROWDAQ*
Qiang Ning, Hao Wu, Pradeep Dasigi, Dheeru Dua, Matt Gardner, Robert L. Logan IV, **Ana Marasović**, Zhen Nie • EMNLP System Demonstrations Track 2020
- [24] *Don't Stop Pretraining: Adapt Language Models to Domains and Tasks*
Suchin Gururangan, **Ana Marasović**, Swabha Swayamdipta, Kyle Lo, Iz Beltagy, Doug Downey, Noah A. Smith • ACL 2020 • **Honorable Mention for Best Paper Award**
- [25] *Quoref: A Reading Comprehension Dataset with Questions Requiring Coreferential Reasoning*
Pradeep Dasigi, Nelson F. Liu, **Ana Marasović**, Noah A. Smith, Matt Gardner • EMNLP 2019
- [26] *SRL4ORL: Improving Opinion Role Labeling Using Multi-Task Learning with Semantic Role Labeling*
Ana Marasović, Anette Frank • NAACL 2018
- [27] *A Mention-Ranking Model for Abstract Anaphora Resolution*
Ana Marasović, Leo Born, Juri Opitz, Anette Frank • EMNLP 2017
- [28] *Modal Sense Classification At Large: Paraphrase-Driven Sense Projection, Semantically Enriched Classification Models and Cross-Genre Evaluations*
Ana Marasović, Mengfei Zhou, Alexis Palmer, Anette Frank • Linguistic Issues in Language Technology (LiLT) – Modality: Logic, Semantics, Annotation, and Machine Learning 14 (3), 2016

Workshop Papers & Other Published Work

- [1] *Teaching People LLM's Errors and Getting it Right*
Nathan Stringham, Fateme Hashemi Chaleshtori, Xinyuan Yan, Zhichao Xu, Bei Wang, **Ana Marasović** • ACL 2026 Workshop on Trustworthy NLP
- [2] *Computer Use Survey: A Visual Survey of Computer Use Agents*
Kenneth Marino, Md Farhan Ishmam, **Ana Marasović** • ICLR 2026 Blogpost Track
- [3] *ConeCut: Scalable Removal of Preference Redundancy*
Purbid Bambroo, Daniel S. Brown, **Ana Marasović** • CoLM 2025 Workshop on Socially Responsible Language Modelling Research (SoLaR)
- [4] *On Evaluating Explanation Utility for Human-AI Decision-Making in NLP*
Fateme Hashemi Chaleshtori, Atreya Ghosal, **Ana Marasović** • The workshop “XAI in Action: Past, Present, and Future Applications” co-located with NeurIPS 2023
- [5] *How Much Consistency Is Your Accuracy Worth?*
Jacob K. Johnson, **Ana Marasović** • The Workshop on Analyzing and Interpreting Neural Networks for NLP (BlackboxNLP) co-located with EMNLP 2023
- [6] *What's Important in a Text? An Extensive Evaluation of Linguistic Annotations for Summarization*
Markus Zopf, Teresa Botschen, Tobias Falke, Benjamin Heinzerling, **Ana Marasović**, Todor Mihaylov, S. Avinesh P.V., Eneldo Loza Mencía, J. Fürnkranz, A. Frank • SNAMS 2018
- [7] *Multilingual Modal Sense Classification using a Convolutional Neural Network*
Ana Marasović, Anette Frank • Workshop on Representation Learning for NLP (Repl4NLP) 2016

Other

- [1] *NLP's Generalization Problem, and How Researchers Are Tackling It*
Ana Marasović • The Gradient 2018

Talks

Keynote Talks

If You Want Reasoning, Look Inside

- The 1st Workshop on the Application of LLM Explainability to Reasoning & Planning @ CoLM 2025 Oct 2025
- The 10th Workshop on Representation Learning for NLP (RepL4NLP 2025) @ NAACL 2025 May 2025

Challenges in Fostering (Dis)Trust in AI

Oct 2023

- Workshop on “Trust Perspectives in Machine Learning, Law, and Public Policy” organized by the Institute for Data, Econometrics, Algorithms, and Learning; located at Northwestern University

Contrastive Explanations of NLP Models

Nov 2021

- EMNLP Workshop on “Analyzing and Interpreting Neural Networks for NLP” (BlackBoxNLP)

Resolving Abstract Anaphors in Discourse: Uphill Battles with Neural Ranking Models and Automatic Data Extraction

Jun 2018

- NAACL Workshop on Computational Models of Reference, Anaphora, and Coreference (CRAC)

Invited Talks

TBD

Nov 2026

- University of Colorado Boulder, NLP Retreat

Faithfulness of LLM Reasoning & Its Emerging Questions

Feb 2026

- University of Arizona, CogSci Colloquium

Explainability and Trust

Apr 2024

- Bellairs Invitational Workshop on Contemporary, Foreseeable and Catastrophic Risks of LLMs; Barbados

Challenges in Fostering (Dis)Trust in AI

Apr 2024

- University of Pennsylvania
- University of Pittsburgh

“We propose <insert an explainability method> to increase human trust in AI”: What does this mean and how to go about it?

Apr 2023

- University of Maryland

Formalizing Trust in AI

May 2022

- NSF AI Institute for Research on Trustworthy AI in Weather, Climate, and Coastal Oceanography (AI2ES)

Self-Explaining for Intuitive Interaction with AI

- Carnegie Mellon University Apr 2022
- NEC Labs Europe Apr 2022
- University of Lisbon – Unbabel Seminar Apr 2022
- UC Irvine Mar 2022
- Johns Hopkins University Feb 2022
- Brown University Feb 2022
- University of Chicago Feb 2022
- Ohio State University Feb 2022
- Purdue University, CS Feb 2022
- University of Utah Feb 2022

On Reducing Human-Authored Free-Text Explanations for Training

Dec 2021

- University College London

Explanation Selection Through The Lens of Free-Text & Contrastive Explanations

- USC Machine Learning Symposium Dec 2021
- Georgia Tech NLP Seminar Nov 2021

• University of Sheffield Seminar	Nov 2021
• Purdue University, ECE	Nov 2021
Explaining with Natural Language	Oct 2021
• University of Edinburgh	
Reflecting on Interpretability Desiderata with Visual Commonsense Reasoning	Nov 2019
• University of Washington, Department of Linguistics Seminar	
Teaching Machines Language Understanding	Apr 2019
• Allen School Women’s Research Day, University of Washington	
Resolving Abstract Anaphors in Discourse: Uphill Battles with Neural Ranking Models and Automatic Data Extraction	
• Allen Institute for AI	Oct 2018
• Potsdam University, Applied Computational Linguistics	Jun 2018
• Bochum University, Sprachwissenschaftliches Institut	May 2018

Media Coverage

Podcast Guest Appearances

- WiAIR Podcast: “Can We Trust AI Explanations? Dr. Ana Marasović on AI Trustworthiness, Explainability & Faithfulness” ([link](#))

Research Coverage in Media

- Psychology Today: “Do AI Models Like GPT Really Get the Joke?” ([link](#))
- The Debrief: Comments by Steve Callaghan (Family Guy writer/producer) and Fred Tatasciore (voice-over actor) ([link](#))
- Cornell Chronicle: “That’s funny – but AI models don’t get the joke” ([link](#))
- Wired: “The Efforts to Make Text-Based AI Less Racist and Terrible” ([link](#))
- Wired: “Review: We Put ChatGPT, Bing Chat, and Bard to the Test” ([link](#))
- Unite.AI: “Minority Voices ‘Filtered’ Out of Google Natural Language Processing Models” ([link](#))

Teaching

CS 5966/6966: Interpretability of LLMs	Spring 2026
CS 5340/6340: Natural Language Processing	Fall 2026, Fall 2025, Spring/Fall 2024
CS 7930: Mechanistic Interpretability Seminar	Spring 2025
CS 7931: Advances in Generative AI Seminar	Fall 2024
CS 5966/6966: Local Explanations for Deep Learning Models	Fall 2023, 2022
CS 5140/6140: Data Mining	Spring 2023

Guest Lectures & Conference Tutorials & Presentations at Local Events

If You Want Reasoning, Look Inside

- UCLA, CS 269 AI Safety & Mechanistic Interpretability, *Host*: Saadia Gabriel Nov 2025
- University of Utah, CS 3020 Research Forum, *Host*: Erik Brunvand Nov 2025

NAACL’24 Tutorial “Explanation in the Era of Large Language Models”

June 2024

Ingredients of Generative AI

- University of Utah, Research Challenges in Computing Seminar, *Host*: Saday Sadayappan Nov 2023
- University of Utah, DELPHI Data Science Symposium Oct 2023
- University of Utah, The Osher Lifelong Learning Institute Lunch and Learn Oct 2023

NLP with (Almost) No Training Data

- University of Utah, MATH 4100 / COMP 5360, *Host*: Kate Isaacs March 2023
- University of Utah, CS 5340/6340, *Host*: Ellen Riloff Nov 2022

“AI” That Masters Language Could Reason about Negation

- University of Utah Data Science Seminar Feb 2023

Towards Trustworthy NLP: On Intuitive Interaction and Rigorous Validation

• University of Utah, CS 3020 Research Forum, <i>Host</i> : Kobus Van Der Merwe	Sep 2022
A Quick Tour of NLP Explainability	
• University of Washington, <i>Host</i> : Yulia Tsvetkov	May 2022
Detecting Artifacts with Local Explanations	
• University of British Columbia, <i>Host</i> : Vered Shwartz	March 2022
• University of Washington, <i>Host</i> : Zaid Harchaoui	March 2022
Methods for Analyzing (Current) NLP Models	
• University of Washington, <i>Host</i> : Noah A. Smith	March 2020

Mentoring

Utah PhD Students

• Fateme Hashemi Chaleshtori	Jan 2023 – present
• Alex Gill	Aug 2025 – present
• Jayanta Sadhu	Oct 2025 – present
• Shruti Ghela	May 2026 – present
• Andrei Kopanev	Jul 2026 – present

Utah MS Researchers

• Neeraja Vasa	Aug 2026 – present
• Thomas Kauffman	May 2026 – present
• Xuyen Nguyen	Jan 2026 – present
• Neha Bhat	Jan 2026 – present
• Nathan Stringham	Aug 2023 – Apr 2026
• Braden Fiedel	Aug 2025 – Aug 2026
• Michael Clemens → New Jersey Institute of Technology PhD	Jul 2023 – Oct 2025
• Purbid Bambroo → Noetica	Jan – May 2024; Sep 2024 – Aug 2025
• Alex Gill → University of Utah PhD	Jan 2024 – Aug 2025
• Rishanth Rajendhran → University of Maryland PhD	Nov 2022 – May 2024
• Adwait Shinganwade	Aug 2024 – Dec 2024
• Yash Lele	Aug 2024 – Dec 2024
• Mugdha Abhyankar	Jan 2023 – May 2023

Utah BS Researchers

• Lucas Pearch	Jan 2026 – present
• Parker Henry DeYoung	Jan 2026 – May 2026
• Matthew Lee	Jan 2026 – May 2026
• Braden Fiedel	Aug 2024 – Aug 2025
• Casey Miller	Aug 2024 – present
• Kishan Thambu	Aug 2024 – present
• Byron Liu (co-mentoring with the PhD student Maitrey Mehta)	Jul 2023 – May 2024
• Jordan Tan	Jul 2023 – Aug 2023
• Dana Ahn	Mar 2023 – Aug 2023
• Jacob Johnson → Utah PhD → Weber State University MS	Nov 2022 – May 2024

UW BS/MS Researchers

• Kaiser Sun with Noah A. Smith → META Residency → JHU PhD	Mar 2020 – present
• Millicent Li with Noah A. Smith → META Residency → Northeastern PhD	Mar 2020 – Aug 2021

AI2 Predoctoral Researchers

• Alexis Ross with Matthew E. Peters → MIT PhD	July 2020 – June 2022
--	-----------------------

AI2 Interns (PhD students)

• Abhilasha Ravichander – Carnegie Mellon University – with Matt Gardner	Jun 2021 – June 2022
• Shruti Palaskar – Carnegie Mellon University	May 2021 – June 2022
• Mimansa Jaiswal – University of Michigan	Sep 2021 – Dec 2021
• Alexander M. Hoyle – University of Maryland – with Noah A. Smith	Jun 2020 – Nov 2020

- Sarah Wiegrefe – Georgia Tech – **AI2 Intern of the Year** – with Noah A. Smith May 2020 – Oct 2020
- Zhanming Allan Jie – Singapore Uni. of Technology & Design – with Pradeep Dasigi Jun 2019 – Aug 2019

Service

Department & University Service

- Faculty Hiring Committee Jan 2026 – Apr 2026
- Faculty Upskilling in AI (U-AI) instructor and consultant Fall 2024, Spring 2025
- AI Track Director Jul 2023 – present
- One Utah RAI Healthcare Working Group Feb 2024 – Jul 2024
- Faculty Hiring Committee Nov 2023 – May 2024
- PhD (AI) Admissions Committee Dec 2022 – Jan 2023

PhD Committee Member

- Oliver Bentham
- Laura Vietz
- Ashim Gupta
- Maitrey Mehta
- Jakob Johnson
- Mattia Medina-Grespan
- Zhichao Xu
- Peter Hase (UNC Chapel Hill)
- Zhenduo Wang
- Marcos Treviso (University of Lisbon)
- Prince Osei Aboagye
- Tianyu Jiang

Founding Member

- International Consortium for Interpretable AI Sep 2026 – present

Board Member

- NAACL Executive Board Jan 2026 – present

Co-organizer

- TeachingNLP: Workshop on Teaching Natural Language Processing @ ACL 2024
- SustainNLP: Workshop on Simple and Efficient Natural Language Processing @ EMNLP 2021
- AllenNLP 2021 Hackathon for Minoritized Communities
- Host of the NLP Highlights Podcast from 2021-2022
- AIPHES Scientific Workshop 2017

Best Paper Award Committee

- EMNLP 2025
- NAACL 2024

Tutorial Chair

- EMNLP 2026

Senior Area Chair

- ACL 2024
- ACL 2023 (**Best Senior Area Chair**)
- EMNLP 2022

Area Chair

- CoLM 2026

Monograph Reviewer

- Reviewer for “Synthesis Lectures in Human Language Technologies: Explainable Natural Language Processing” in 2021

Journal Reviewer

- TACL 2017 (secondary reviewer), 2021; *in 2021 TACL had an impact factor of 9.194*

Conference Reviewer

- ARR May 2025, ARR July 2025, ARR May 2026
- NAACL 2024
- EMNLP 2018, 2019 (**outstanding reviewer**), 2020 (**outstanding reviewer**), 2021, 2024
- ACL 2017 (secondary reviewer), 2019, 2020, 2021
- EACL 2021, 2023
- ICML 2020
- AAAI 2020

Workshop Reviewer

- Representation Learning for NLP (Repl4NLP) @ ACL 2018, ACL 2019
- Deep Learning Approaches for Low-Resource Natural Language Processing (DeepLo) @ EMNLP 2019
- Learning from Limited Labeled Data (LLD) Workshop @ ICLR 2019
- Evaluation and Comparison of NLP Systems (EvalNLP) @ EMNLP 2020
- Insights from Negative Results in NLP @ EMNLP 2020
- Computational Approaches to Discourse (CODI) @ EMNLP 2020
- Domain Adaptation for NLP (Adapt-NLP) @ EACL 2021

Panelist

- Hinckley forum panelist on the theme of “AI in policy and society” in Oct 2026
- Moderator of the panel “AI in Pop Culture vs. Reality” at the Leonardo in Nov 2024
- Panelist at the Utah Digital Health Initiative “Generative AI in Health Sciences Research” panel in Nov 2023
- Panelist at the SCI’s 30th Anniversary event in April 2024
- Panelist at a career panel at 2023 EACL Student Research Workshop (SRW)
- Panelist for the screening of the movie “After Yang” for the STEM Education Film Series by Park City Film in Sep 2022
- Panelist at ACL Mentorship: “Building NLP datasets” in Aug 2022
- Panelist at SoCal ML & NLP Symposium 2021 (Networking and career paths for under-represented minorities)
- Panelist at undergraduate panel session @ ACL 2020

Moderator at Conferences

- Moderator of a meetup on “Interpretability” @ ACL 2021
- Moderator of a meetup on “Whether to do a PhD; and how to apply for and choose a PhD program” @ EMNLP 2020
- Moderator of a meetup on “Interpretability and Analysis of Models for NLP” at EMNLP 2020
- Moderator of a roundtable on “Natural Language Generation” at WiNLP @ ACL 2020

Mentoring at Conferences

- Mentor at WiNLP @ ACL 2020
- “Big sibling” at EMNLP 2019, NAACL 2019

Other Service

- Leonardo AI Advisory Committee Member, Jun 2024–2025
- Co-Lead of the AI2 Employee Resource Group (ERG) for Women and Gender Minorities in 2022
- Member of the AI2 Diversity, Equity and Inclusion Committee in 2020 and 2021
- Reader of 2020 PhD applications at the University of Washington

References

Noah A. Smith, nasmith@cs.washington.edu

Yejin Choi, yejinc@stanford.edu

Anette Frank, frank@cl.uni-heidelberg.de

Doug Downey, dougdowney@allenai.org